

digit-detection-using-cnn

May 4, 2023

Importing data from Kaggle

```
[1]: ! pip install -q kaggle
      from google.colab import files
      files.upload()
      ! mkdir ~/.kaggle
      ! cp kaggle.json ~/.kaggle/
      ! chmod 600 ~/.kaggle/kaggle.json
```

<IPython.core.display.HTML object>

Saving kaggle.json to kaggle (2).json

mkdir: cannot create directory '/root/.kaggle': File exists

```
[2]: !kaggle datasets download -d devanshusingh/
      ↪hand-written-digits-for-object-detection
```

hand-written-digits-for-object-detection.zip: Skipping, found more recently modified local copy (use --force to force download)

```
[ ]: !unzip hand-written-digits-for-object-detection.zip
```

Importing Dependencies

```
[4]: import os
      import cv2 as cv
      import keras
      import numpy as np
      import pandas as pd
      import seaborn as sns
      import matplotlib.pyplot as plt
      import matplotlib.patches as patches
      from sklearn.model_selection import train_test_split
      from keras.layers import Conv2D, MaxPool2D, Dense, GlobalAveragePooling2D, Input
      from keras.models import Sequential, Model
      from sklearn.metrics import mean_absolute_error
```

```
[5]: df = pd.read_csv('/content/train_label.csv')
```

```
[6]: df.head()
```

```
[6]:  image_path  label  x_mark  y_mark  width  height
0      0.png      5    23.0    0.0    51.0    28.0
1      1.png      0    43.0    34.0    71.0    62.0
2      2.png      4    43.0    23.0    71.0    51.0
3      3.png      1     3.0    27.0    31.0    55.0
4      4.png      9    32.0    12.0    60.0    40.0
```

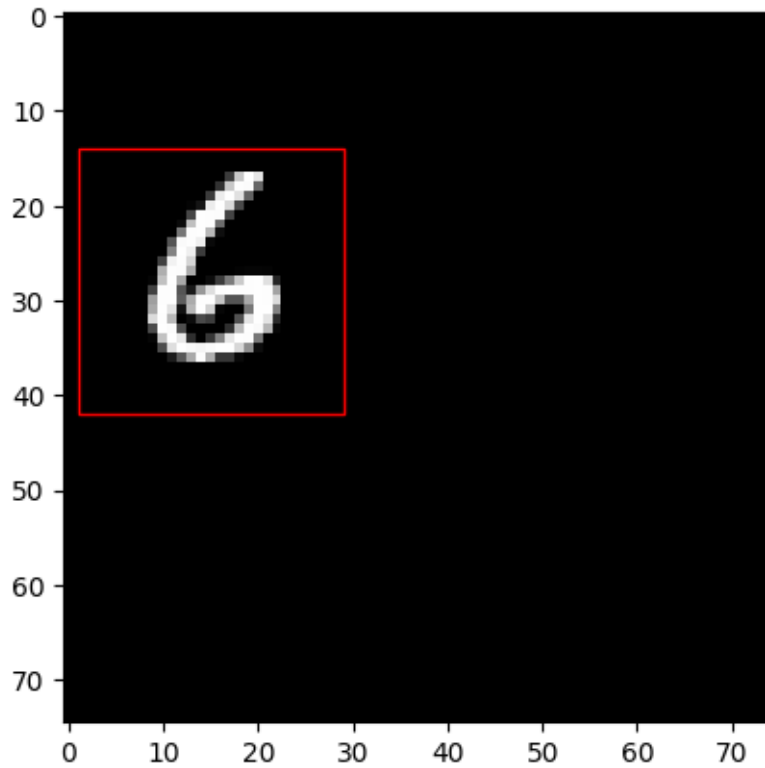
```
[8]: data = {} # appending the bbox
      ↪coordinates into the data dictionary
for idx, names in enumerate(df["image_path"]):
    if names not in data:
        data["train/train/" + names] = []
        x1 = int(df["x_mark"][idx])
        y1 = int(df["y_mark"][idx])
        x2 = int(df["width"][idx])
        y2 = int(df["height"][idx])
        data["train/train/" + names].append(x1)
        data["train/train/" + names].append(y1)
        data["train/train/" + names].append(x2)
        data["train/train/" + names].append(y2)
```

```
[9]: images = [] # appending the images into the images array
for keys in data.keys():
    img_arr = cv.imread(keys, cv.IMREAD_GRAYSCALE)
    images.append(img_arr)
```

```
[10]: bbox = []
for bboxes in data.keys():
    bbox.append(data[bboxes])
```

```
[11]: images = np.array(images)
bbox = np.array(bbox)
```

```
[12]: def test_plot(img_number):
        x1 = bbox[img_number][0]
        y1 = bbox[img_number][1]
        width = bbox[img_number][2]
        height = bbox[img_number][3]
        plt.subplot(1,1,1)
        plt.imshow(images[img_number], cmap='gray')
        plt.gca().add_patch(patches.Rectangle((x1, y1), width-x1, height-y1,
        ↪edgecolor='r', facecolor="none", linewidth=1))
    plt.show()
    test_plot(6000)
```



```
[13]: images = np.expand_dims(images, axis=3)      # adding a new axis to the images
      ↪ array
```

```
[14]: images = images/255                        # normalising the images
```

Splitting and Model Building

```
[15]: x_train, x_test, y_train, y_test = train_test_split(images, bbox, test_size=0.
      ↪ 2, random_state=44)
```

```
[16]: model = Sequential()

model.add(Conv2D(32, (3,3), input_shape=(75,75,1), activation="relu"))
model.add(MaxPool2D(2,2))
model.add(Conv2D(64, (3,3), activation="relu"))
model.add(MaxPool2D(2,2))
model.add(Conv2D(128, (3,3), activation="relu"))
model.add(MaxPool2D(2,2))
model.add(GlobalAveragePooling2D())
model.add(Dense(256, activation="relu"))
model.add(Dense(4, activation="relu"))
```

```
[17]: model.summary()
```

```
Model: "sequential"
```

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 73, 73, 32)	320
max_pooling2d (MaxPooling2D)	(None, 36, 36, 32)	0
conv2d_1 (Conv2D)	(None, 34, 34, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 17, 17, 64)	0
conv2d_2 (Conv2D)	(None, 15, 15, 128)	73856
max_pooling2d_2 (MaxPooling2D)	(None, 7, 7, 128)	0
global_average_pooling2d (GlobalAveragePooling2D)	(None, 128)	0
dense (Dense)	(None, 256)	33024
dense_1 (Dense)	(None, 4)	1028

=====
Total params: 126,724
Trainable params: 126,724
Non-trainable params: 0
=====

```
[18]: class CustomCallbacks(keras.callbacks.Callback):  
    def on_epoch_end(self, epoch, logs={}):  
        num = np.random.randint(0, 10000)  
        bbox = model.predict(x_test[num].reshape(1, 75, 75, 1))  
        fig, ax = plt.subplots(1)  
        ax.imshow(x_test[num])  
        x1 = int(bbox[0][0])  
        y1 = int(bbox[0][1])  
        x2 = int(bbox[0][2])  
        y2 = int(bbox[0][3])  
        rect = patches.Rectangle((x1, y1), x2-x1, y2-y1,  
                                linewidth=1, edgecolor='r', facecolor="none")  
        ax.add_patch(rect)
```

```
plt.show()
```

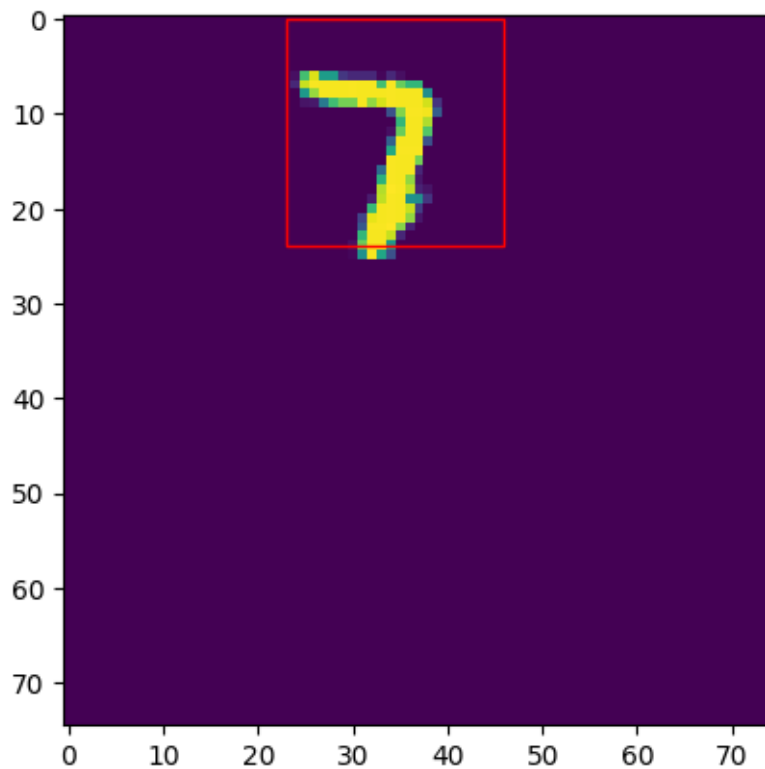
```
[19]: model.compile(optimizer="adam", loss='mae')
```

```
[20]: history = model.fit(x_train, y_train, validation_data=(x_test, y_test),  
    ↪ epochs=30, callbacks=[CustomCallbacks()])
```

Epoch 1/30

1/1 [=====] - 1s 517ms/step

<Figure size 1000x500 with 0 Axes>

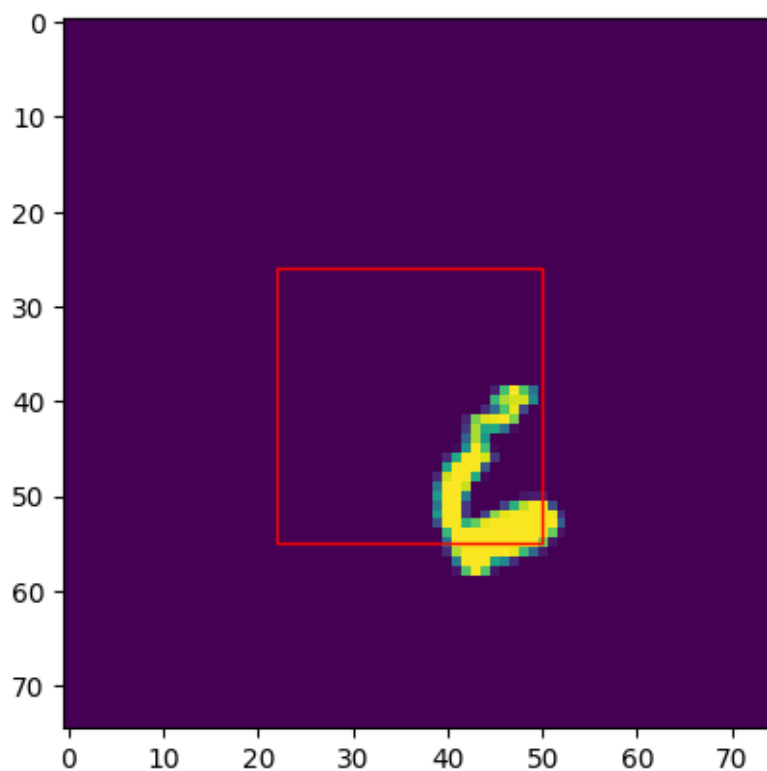


1500/1500 [=====] - 23s 12ms/step - loss: 11.0341 -
val_loss: 6.7351

Epoch 2/30

1/1 [=====] - 0s 19ms/step

<Figure size 1000x500 with 0 Axes>

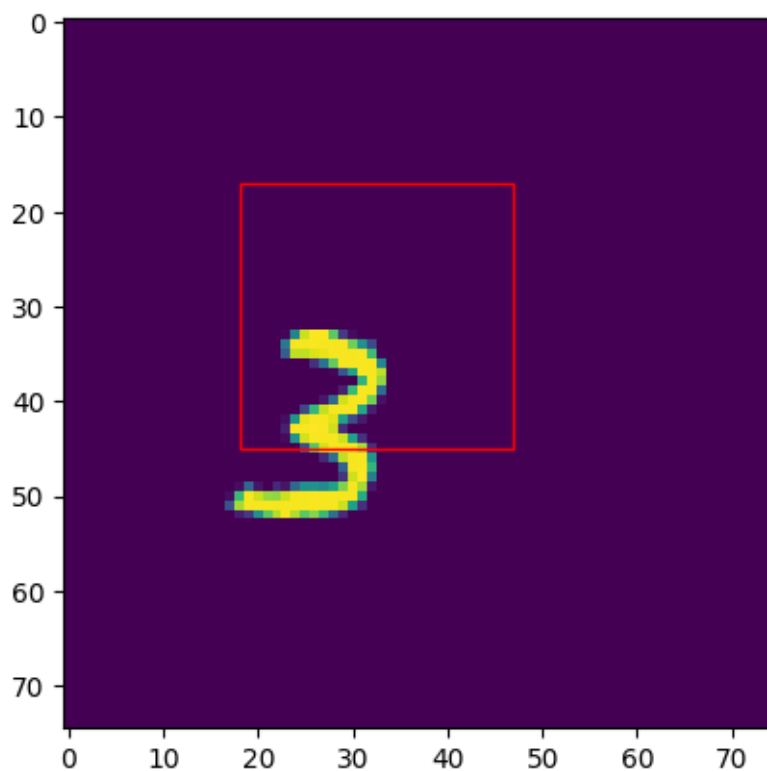


1500/1500 [=====] - 13s 9ms/step - loss: 6.0866 -
val_loss: 5.7306

Epoch 3/30

1/1 [=====] - 0s 31ms/step

<Figure size 1000x500 with 0 Axes>

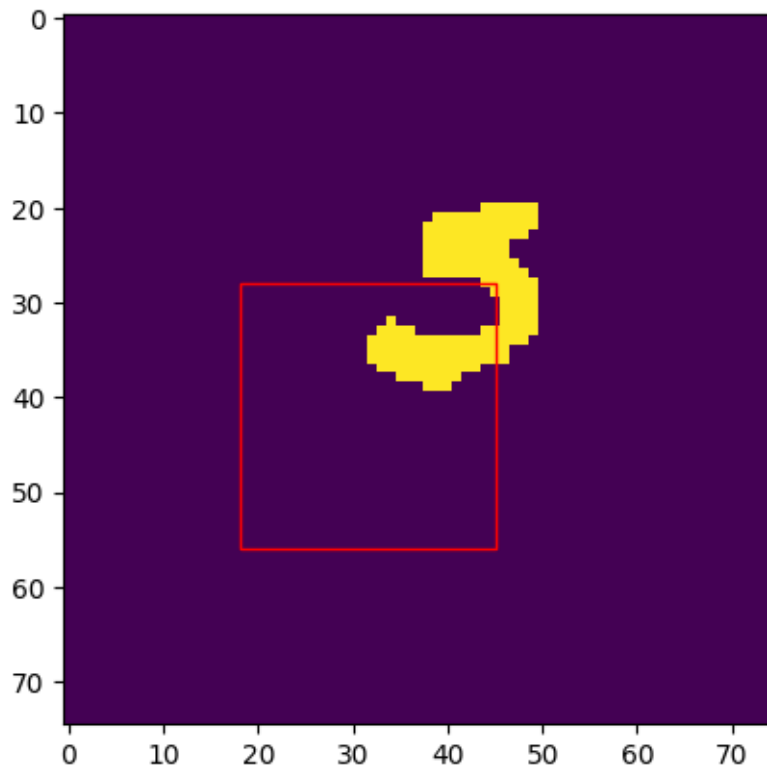


1500/1500 [=====] - 10s 7ms/step - loss: 5.2336 -
val_loss: 5.1354

Epoch 4/30

1/1 [=====] - 0s 20ms/step

<Figure size 1000x500 with 0 Axes>

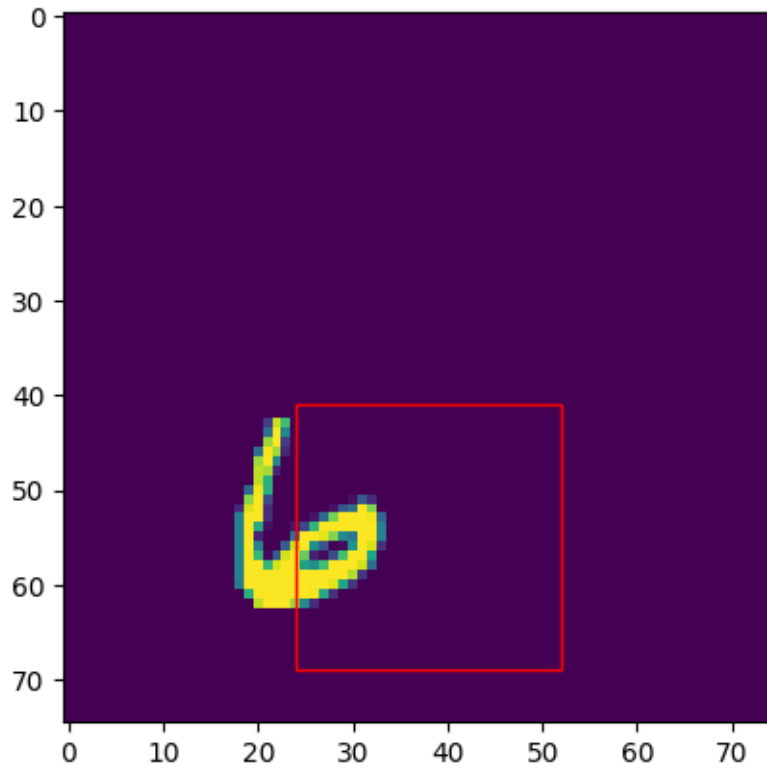


1500/1500 [=====] - 11s 7ms/step - loss: 4.7971 -
val_loss: 5.0530

Epoch 5/30

1/1 [=====] - 0s 18ms/step

<Figure size 1000x500 with 0 Axes>




1500/1500 [=====] - 11s 7ms/step - loss: 4.5436 -
val_loss: 4.9210

Epoch 6/30

1/1 [=====] - 0s 21ms/step

<Figure size 1000x500 with 0 Axes>

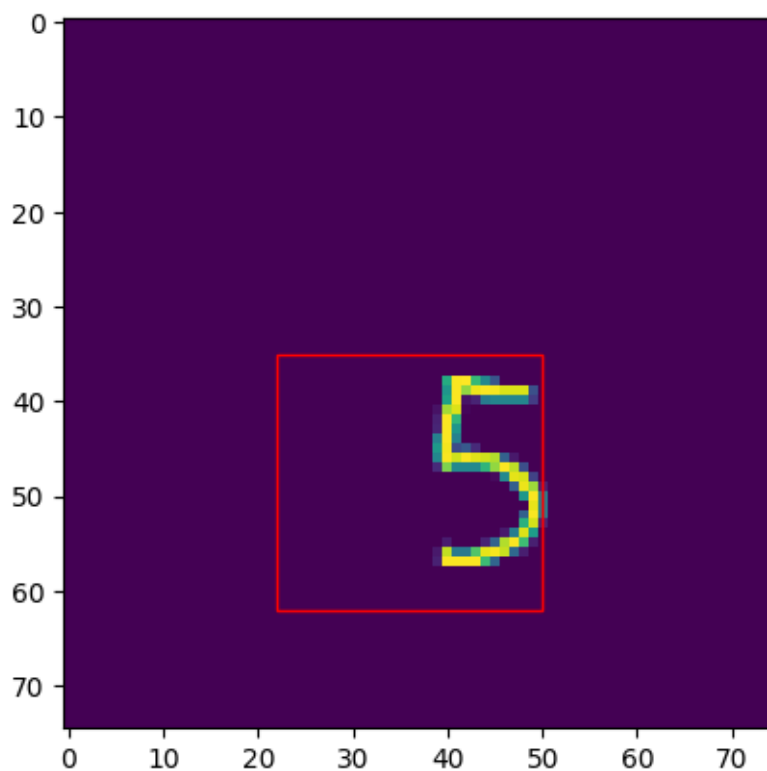


1500/1500 [=====] - 11s 7ms/step - loss: 4.3893 -
val_loss: 4.4408

Epoch 7/30

1/1 [=====] - 0s 26ms/step

<Figure size 1000x500 with 0 Axes>

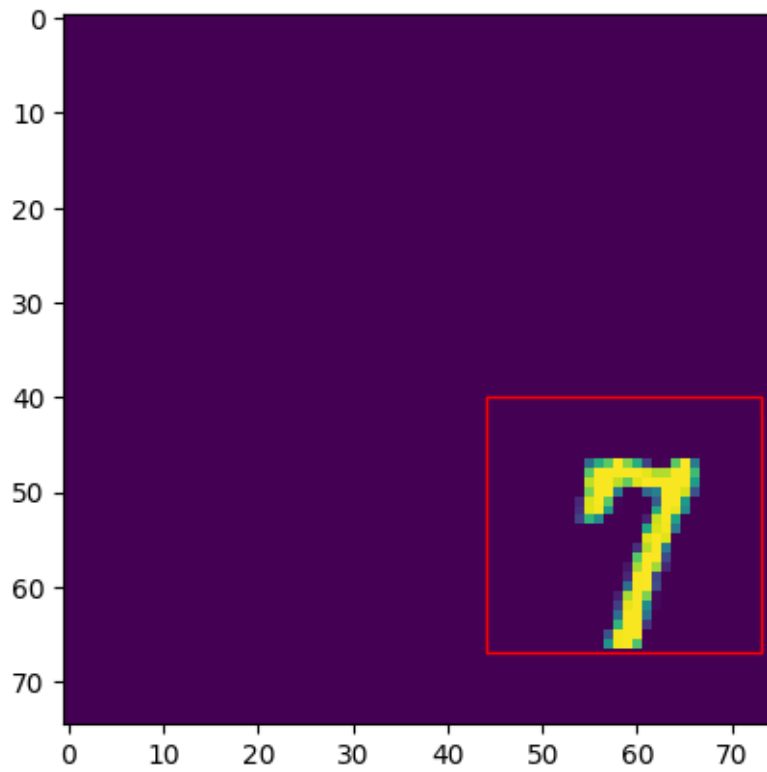


1500/1500 [=====] - 11s 7ms/step - loss: 4.3286 -
val_loss: 4.3952

Epoch 8/30

1/1 [=====] - 0s 19ms/step

<Figure size 1000x500 with 0 Axes>




1500/1500 [=====] - 11s 7ms/step - loss: 4.2556 -
val_loss: 4.3557


Epoch 9/30

1/1 [=====] - 0s 38ms/step

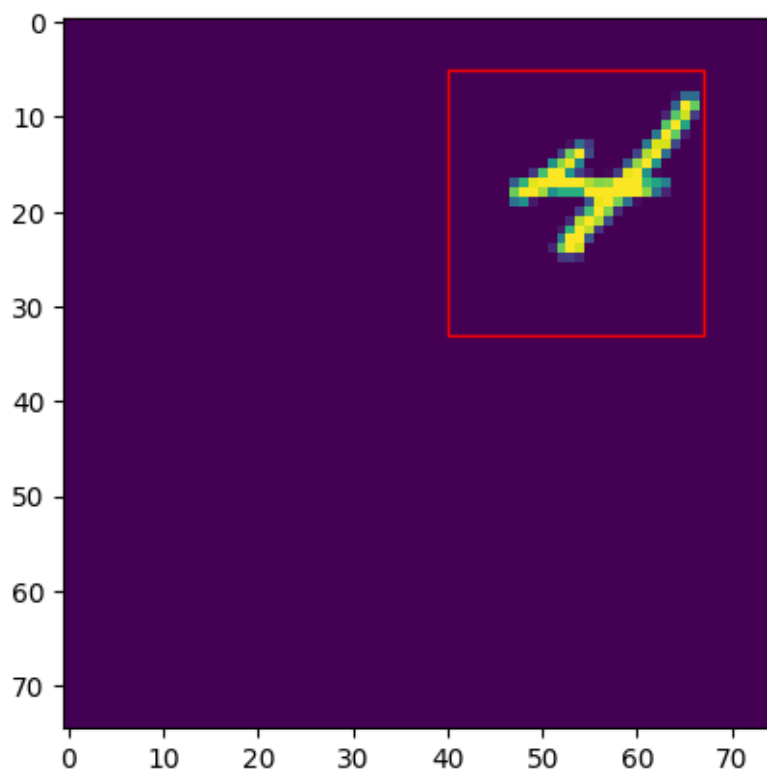
<Figure size 1000x500 with 0 Axes>



```
1500/1500 [=====] - 14s 10ms/step - loss: 4.2034 -  
val_loss: 4.2513  
Epoch 10/30  
1/1 [=====] - 0s 21ms/step  
<Figure size 1000x500 with 0 Axes>
```



```
1500/1500 [=====] - 11s 8ms/step - loss: 4.1453 -  
val_loss: 4.3206  
Epoch 11/30  
1/1 [=====] - 0s 26ms/step  
<Figure size 1000x500 with 0 Axes>
```

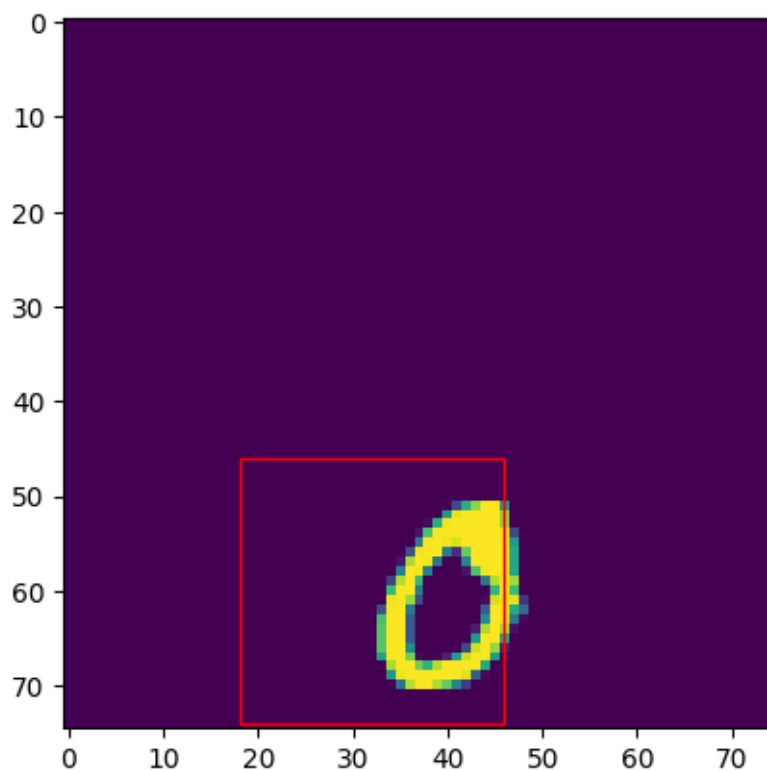


1500/1500 [=====] - 11s 7ms/step - loss: 4.1277 -
val_loss: 4.0900

Epoch 12/30

1/1 [=====] - 0s 29ms/step

<Figure size 1000x500 with 0 Axes>

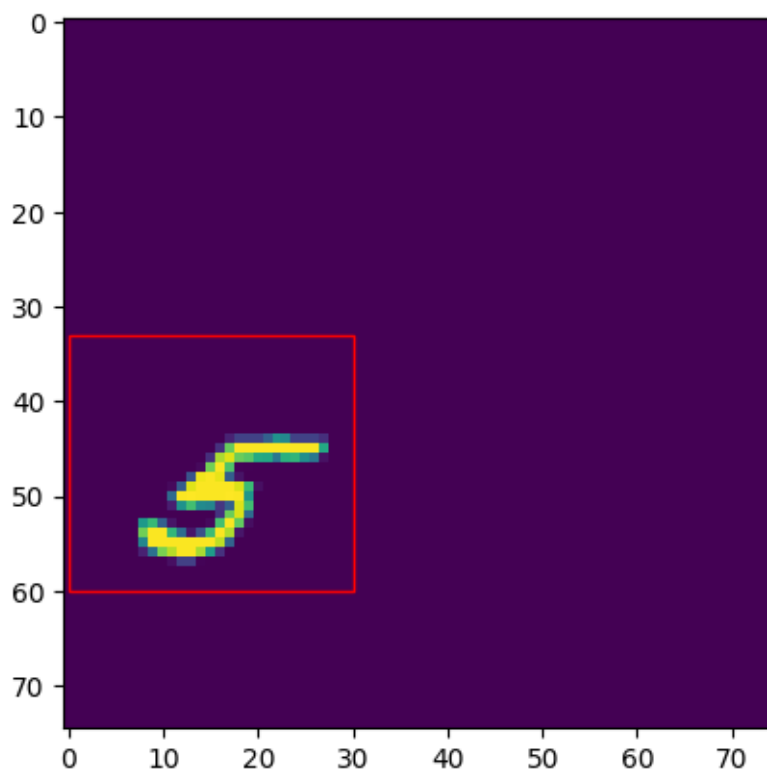


1500/1500 [=====] - 11s 8ms/step - loss: 4.0865 -
val_loss: 4.0703

Epoch 13/30

1/1 [=====] - 0s 21ms/step

<Figure size 1000x500 with 0 Axes>

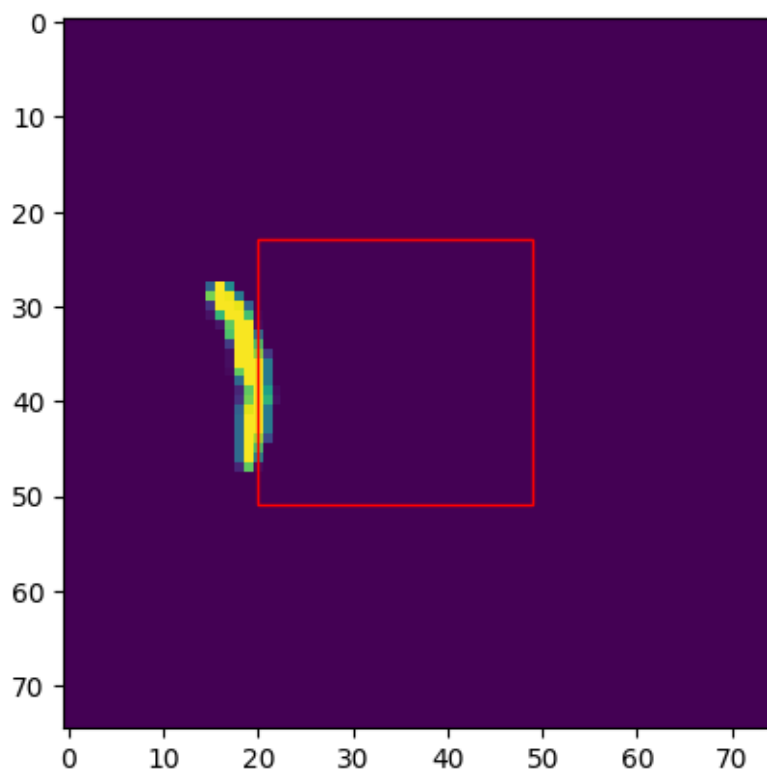


1500/1500 [=====] - 11s 7ms/step - loss: 4.0419 -
val_loss: 4.0991

Epoch 14/30

1/1 [=====] - 0s 18ms/step

<Figure size 1000x500 with 0 Axes>

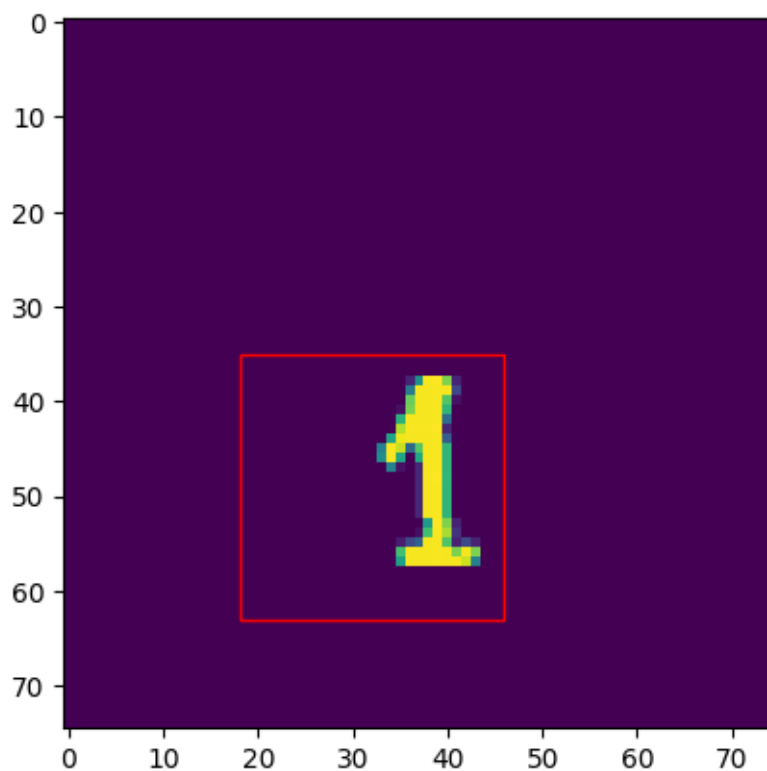


1500/1500 [=====] - 11s 7ms/step - loss: 4.0640 -
val_loss: 3.9957

Epoch 15/30

1/1 [=====] - 0s 22ms/step

<Figure size 1000x500 with 0 Axes>




1500/1500 [=====] - 11s 8ms/step - loss: 4.0109 -
val_loss: 3.9401

Epoch 16/30

1/1 [=====] - 0s 21ms/step

<Figure size 1000x500 with 0 Axes>

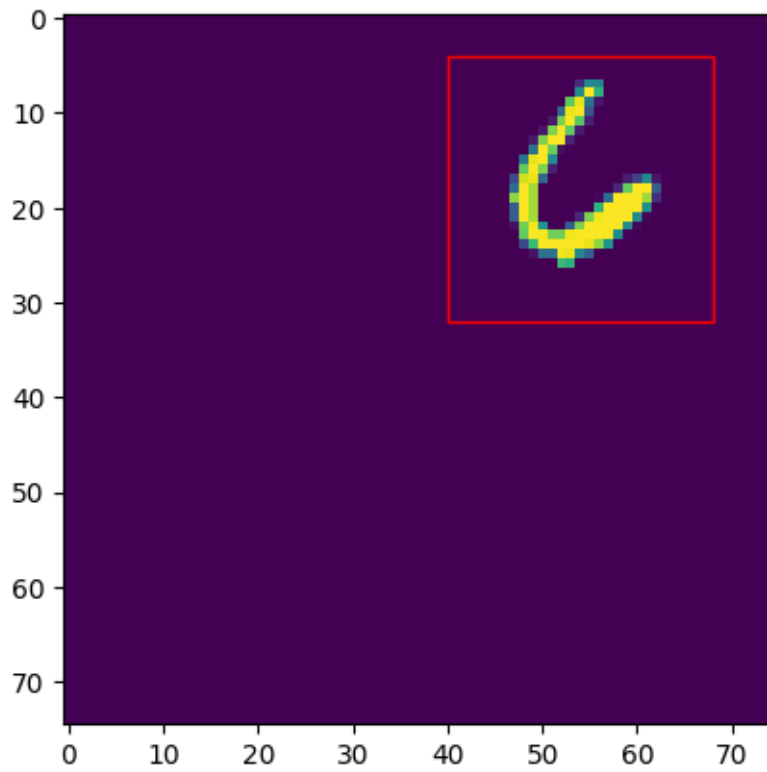


1500/1500 [=====] - 11s 7ms/step - loss: 3.9985 -
val_loss: 4.0390

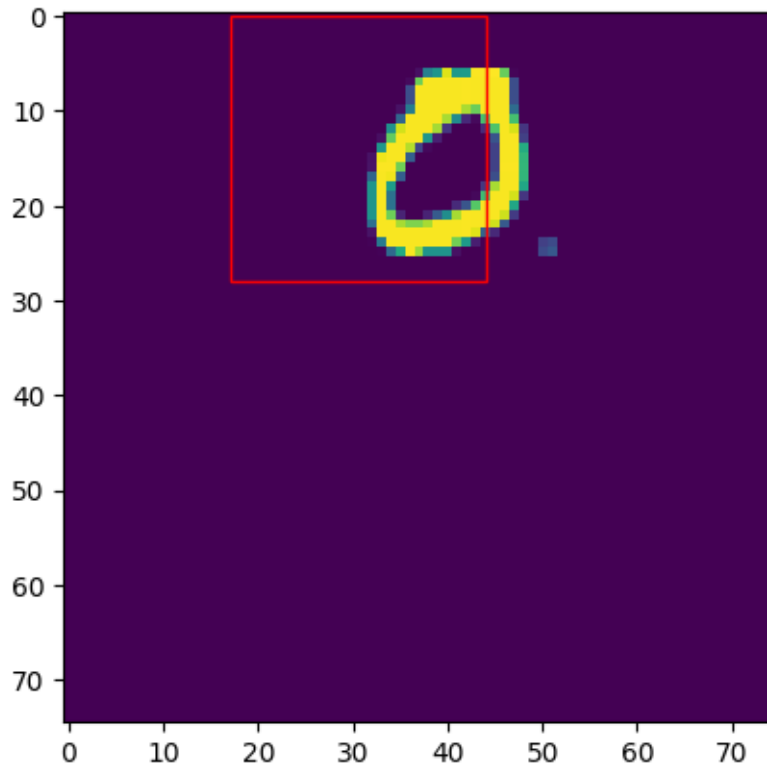
Epoch 17/30

1/1 [=====] - 0s 18ms/step

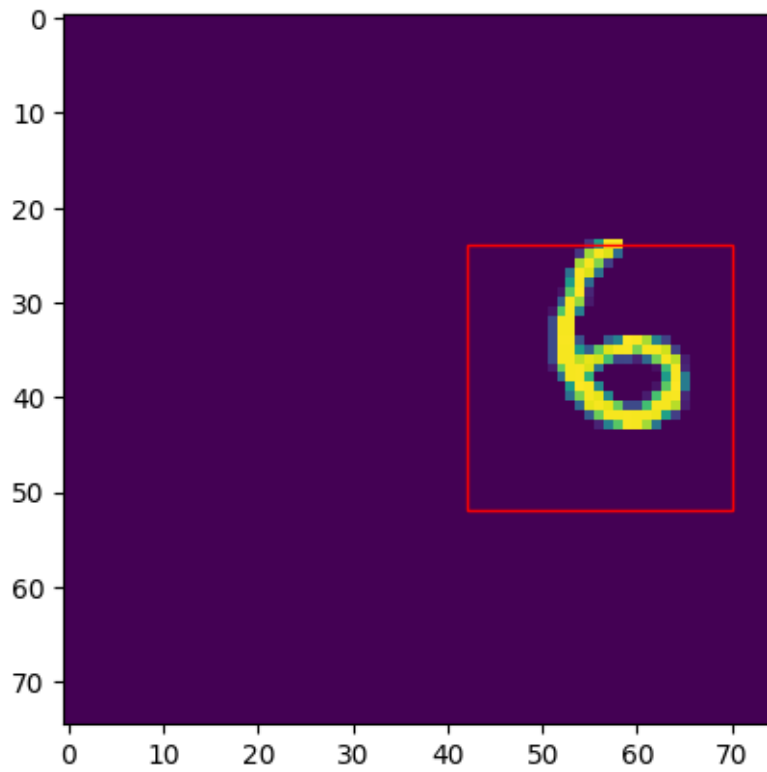
<Figure size 1000x500 with 0 Axes>



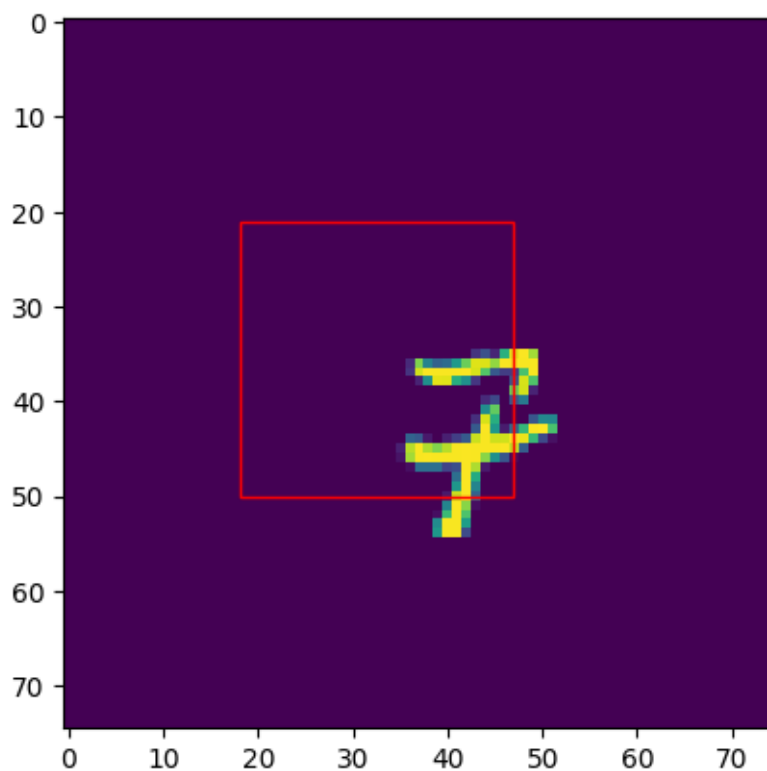
```
1500/1500 [=====] - 11s 7ms/step - loss: 3.9795 -  
val_loss: 3.9221  
Epoch 18/30  
1/1 [=====] - 0s 27ms/step  
<Figure size 1000x500 with 0 Axes>
```



```
1500/1500 [=====] - 11s 7ms/step - loss: 3.9545 -  
val_loss: 3.9476  
Epoch 19/30  
1/1 [=====] - 0s 29ms/step  
<Figure size 1000x500 with 0 Axes>
```



```
1500/1500 [=====] - 12s 8ms/step - loss: 3.9382 -  
val_loss: 3.9146  
Epoch 20/30  
1/1 [=====] - 0s 19ms/step  
<Figure size 1000x500 with 0 Axes>
```

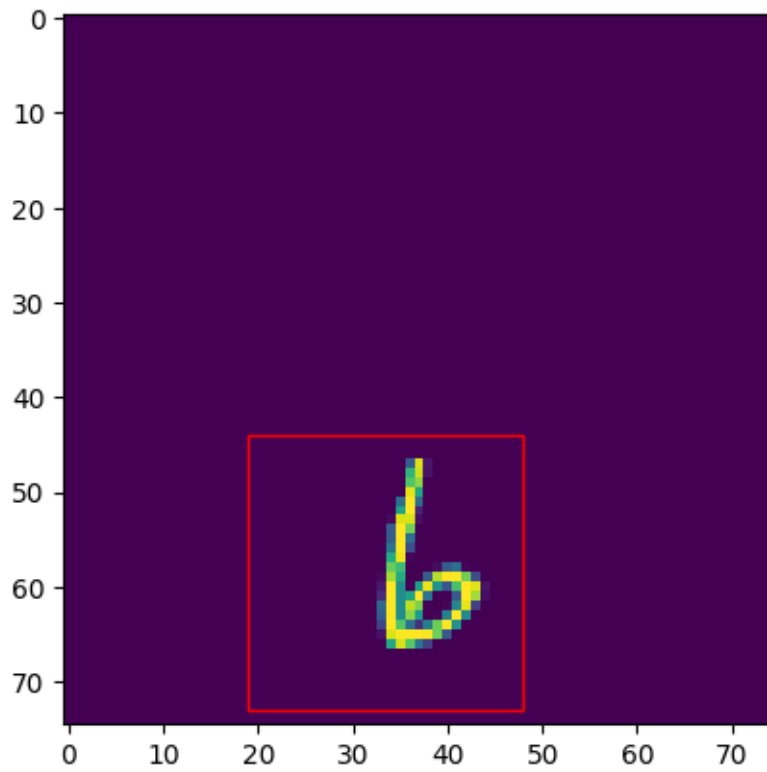


1500/1500 [=====] - 12s 8ms/step - loss: 3.9172 -
val_loss: 3.8593

Epoch 21/30

1/1 [=====] - 0s 21ms/step

<Figure size 1000x500 with 0 Axes>




1500/1500 [=====] - 12s 8ms/step - loss: 3.9241 -
val_loss: 3.9327

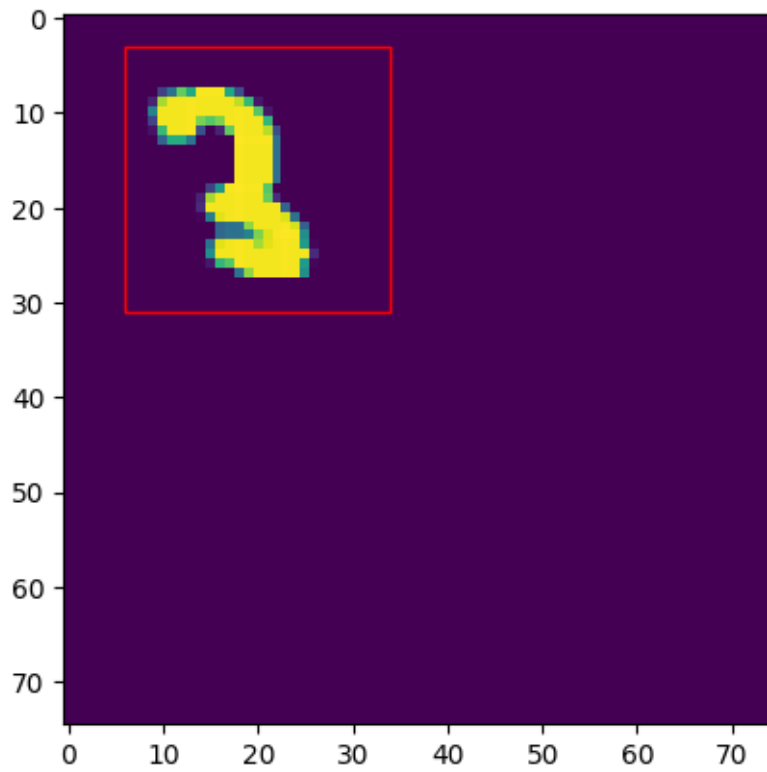
Epoch 22/30

1/1 [=====] - 0s 22ms/step

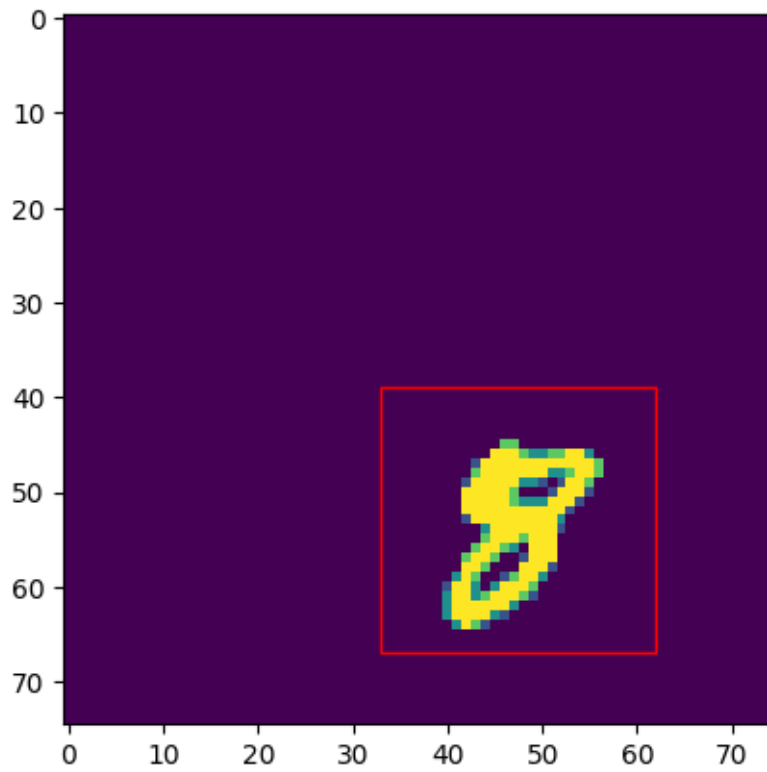
<Figure size 1000x500 with 0 Axes>



```
1500/1500 [=====] - 12s 8ms/step - loss: 3.8779 -  
val_loss: 3.9742  
Epoch 23/30  
1/1 [=====] - 0s 19ms/step  
<Figure size 1000x500 with 0 Axes>
```



```
1500/1500 [=====] - 11s 8ms/step - loss: 3.8836 -  
val_loss: 3.8842  
Epoch 24/30  
1/1 [=====] - 0s 20ms/step  
<Figure size 1000x500 with 0 Axes>
```

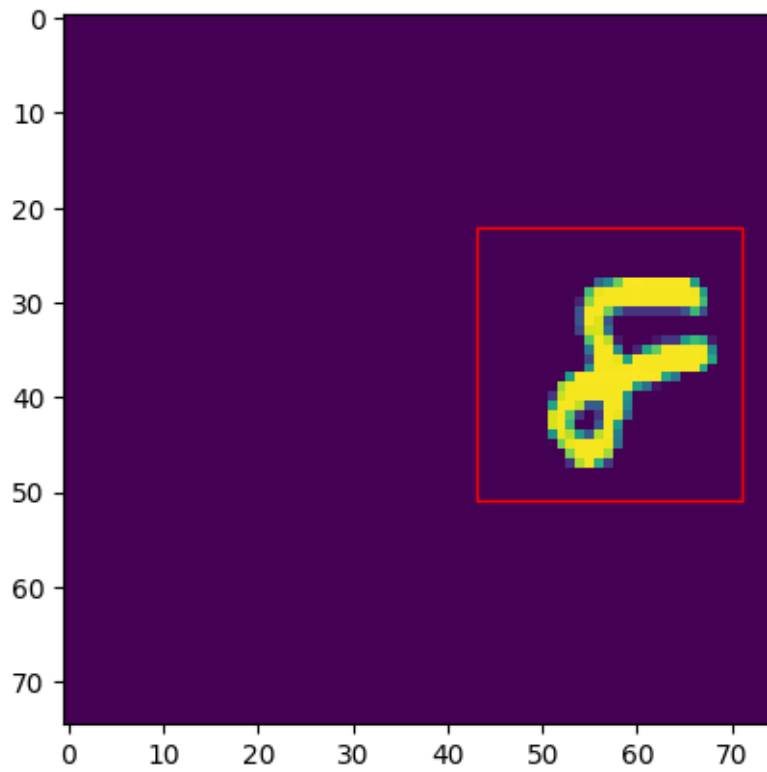


1500/1500 [=====] - 11s 8ms/step - loss: 3.8524 -
val_loss: 3.8593

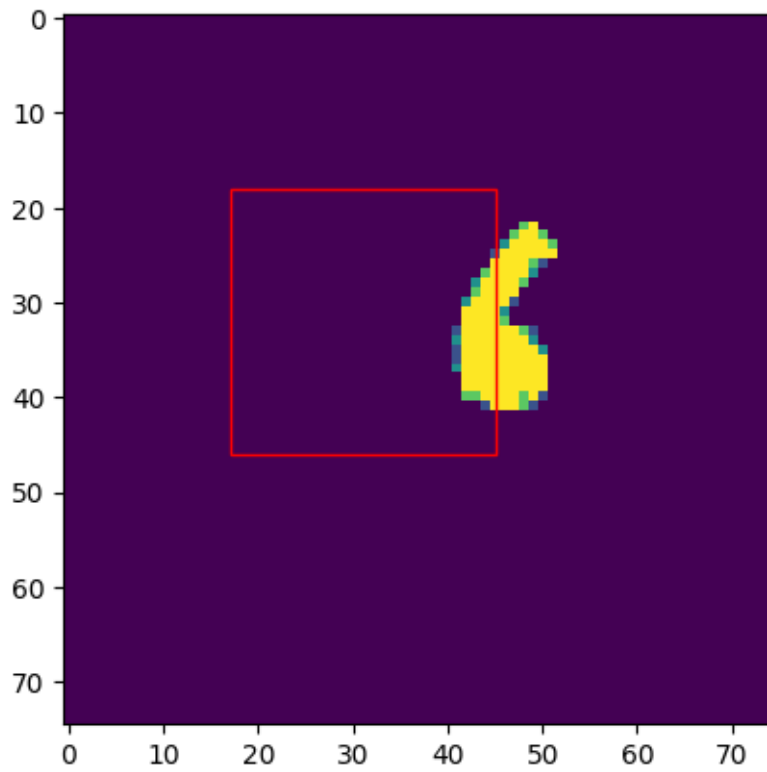
Epoch 25/30

1/1 [=====] - 0s 22ms/step

<Figure size 1000x500 with 0 Axes>



```
1500/1500 [=====] - 11s 8ms/step - loss: 3.8601 -  
val_loss: 3.8985  
Epoch 26/30  
1/1 [=====] - 0s 36ms/step  
<Figure size 1000x500 with 0 Axes>
```

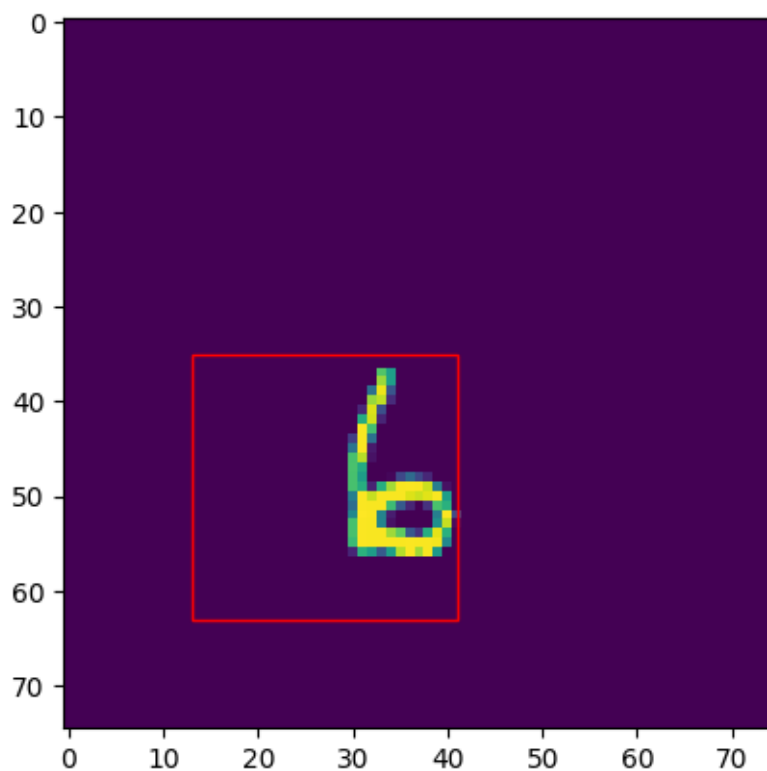


1500/1500 [=====] - 12s 8ms/step - loss: 3.8284 -
val_loss: 3.8247

Epoch 27/30

1/1 [=====] - 0s 42ms/step

<Figure size 1000x500 with 0 Axes>

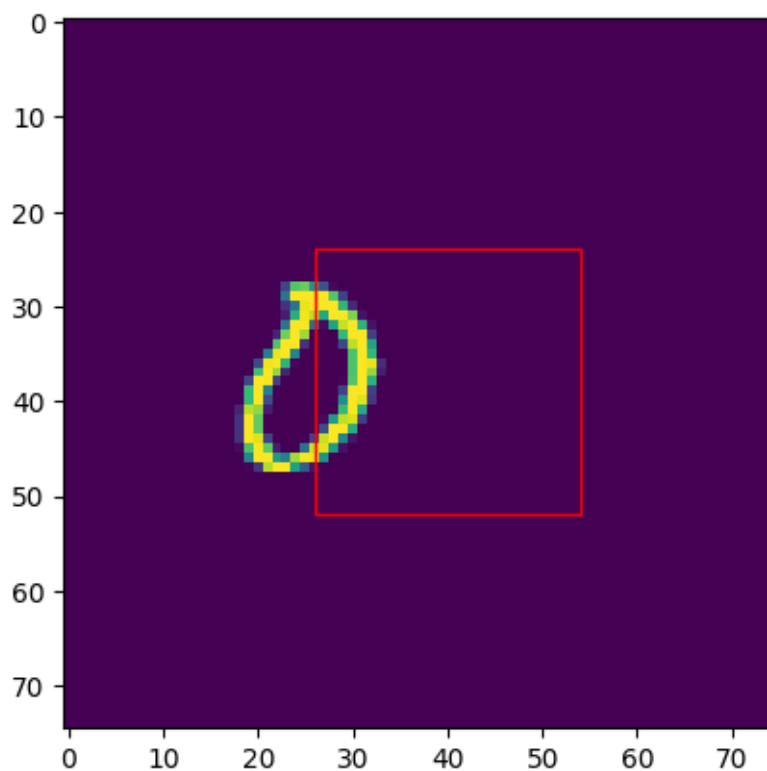


1500/1500 [=====] - 11s 7ms/step - loss: 3.8308 -
val_loss: 3.8820

Epoch 28/30

1/1 [=====] - 0s 19ms/step

<Figure size 1000x500 with 0 Axes>

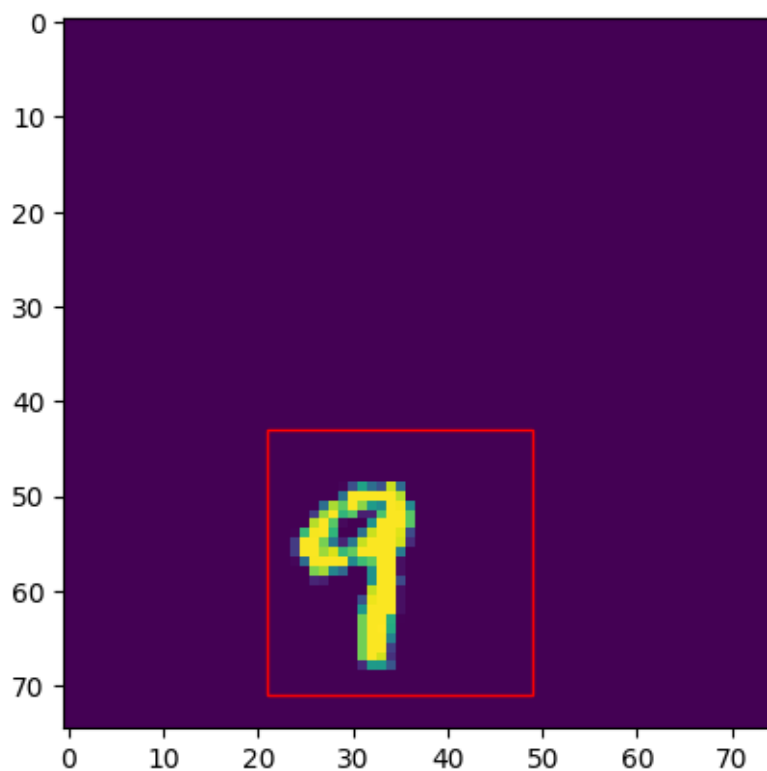


1500/1500 [=====] - 12s 8ms/step - loss: 3.8187 -
val_loss: 4.0059

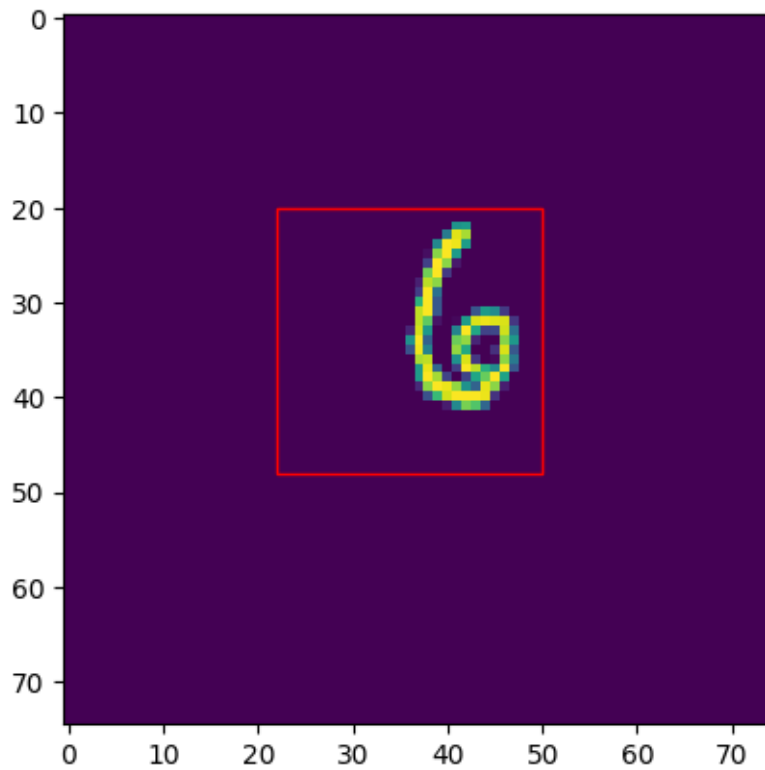
Epoch 29/30

1/1 [=====] - 0s 22ms/step

<Figure size 1000x500 with 0 Axes>

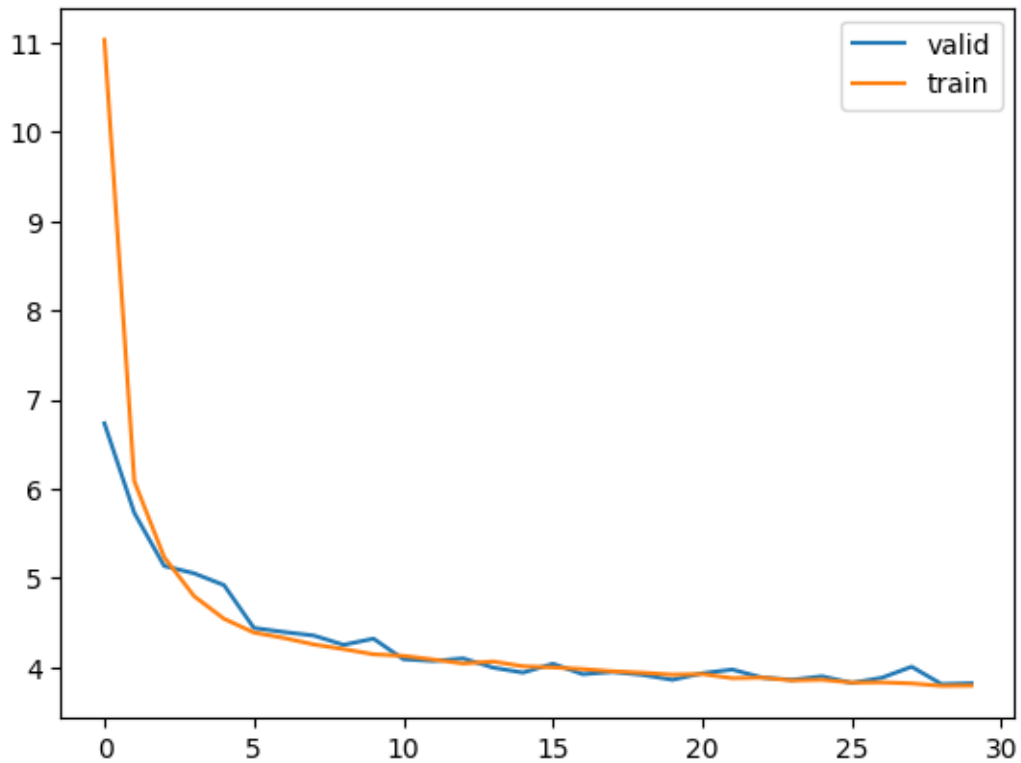


```
1500/1500 [=====] - 12s 8ms/step - loss: 3.7913 -  
val_loss: 3.8108  
Epoch 30/30  
1/1 [=====] - 0s 19ms/step  
<Figure size 1000x500 with 0 Axes>
```



1500/1500 [=====] - 11s 8ms/step - loss: 3.7943 -
val_loss: 3.8207

```
[21]: plt.plot(history.history["val_loss"])  
      plt.plot(history.history["loss"])  
      plt.legend(["valid", "train"])  
      plt.show()
```



```
[22]: def plot_prediction(img_number):
    prediction = model.predict(x_test[img_number].reshape(1, 75, 75, 1))
    x1 = int(prediction[0][0])
    y1 = int(prediction[0][1])
    x2 = int(prediction[0][2])
    y2 = int(prediction[0][3])
    plt.imshow(x_test[img_number])
    plt.gca().add_patch(patches.Rectangle((x1, y1), x2-x1, y2-y1,
    ↪ linewidth=1, edgecolor='r', facecolor="none"))
    plt.show()
plot_prediction(30)
```

1/1 [=====] - 0s 19ms/step

